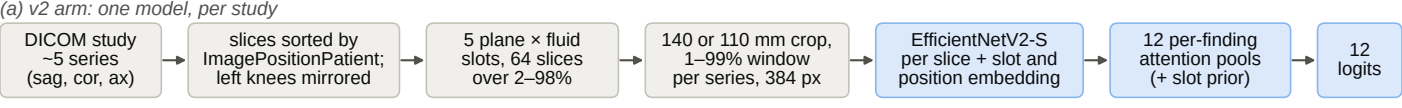
(a) v2 arm: one model, per study
DICOM study ~5 series (sag, cor, ax)
slices sorted by ImagePositionPatient; left knees mirrored
5 plane × fluid slots, 64 slices over 2–98%
140 or 110 mm crop, 1–99% window per series, 384 px
EfficientNetV2-S per slice + slot and position embedding
12 per-finding attention pools (+ slot prior)
12 logits
(b) final ensemble (#13): per-column rank average of three arms
hr arm: 5 folds, weight 3 full frame, 384 px, 24 slices, one attention pool, TTA S = 3
+
v2 arm: 10 models, weight 10 5 folds × {140, 110} mm, TTA S = 3 (panel a)
+
public 2.5D ConvNeXt-T reader: 3 folds, weight 6, no TTA
percentile rank average per column
submission (12 columns)